

Commands: awk and sed.

Jan Trdlička

Department of Computer Systems
Faculty of Information Technology
Czech Technical University in Prague

trdlicka@fit.cvut.cz

November 26, 2018

Contents

- 1 Discussion of homework
- 2 sed
- 3 sed – print/delete
- 4 sed – text replacement
- 5 awk
- 6 awk – patterns
- 7 awk – actions
- 8 awk – functions
- 9 awk – arrays
- 10 awk – examples
- 11 Homework

Discussion of homework

- How to print lines containing the same character at the beginning and at the end of line.

Hint: Use `man -s 5 regexp` and find meaning of symbols `\n`, `\(` and `\)` in the section 2.6 on server `fray1.fit.cvut.cz`.

```
grep '^\(.\).*\1$' /usr/dict/words
```

- How to print all palindromes of length 2, 3, 4 and 5 characters from the file `/usr/dict/words` on server `fray1.fit.cvut.cz`.

```
grep '^\(.\)\1$' /usr/dict/words
```

```
grep '^\(.\)\1$' /usr/dict/words
```

```
grep '^\(.\)\(.\)\2\1$' /usr/dict/words
```

```
grep '^\(.\)\(.\)\1\2\1$' /usr/dict/words
```

- What is the meaning of the `-n` option at command `sed`?
 - It suppresses automatic printing of pattern space (input lines).

```
cat -n /etc/passwd | sed ' '
```

```
cat -n /etc/passwd | sed -n ' '
```

- What is the meaning of the following `sed` commands?
 - `p ...` print the current pattern space (input lines),

```
cat -n /etc/passwd | sed 'p'
```

```
cat -n /etc/passwd | sed -n 'p'
```

- `d ...` delete pattern space (input lines),

```
cat -n /etc/passwd | sed 'd'
```

- `s ...` replace a portion matched regular expression.

```
ls -l / | sed 's/root/ROOT/g'
```

sed – print/delete

- How to print the first 10 lines of the file /etc/passwd by command sed?

```
cat -n /etc/passwd | sed -n '1,10p'
```

```
cat -n /etc/passwd | sed '11,$d'
```

- How to print the last 10 lines of the file /etc/passwd by command sed?

```
cat -n /etc/passwd | \  
sed -n "$((${wc -l </etc/passwd}-9)),\ $p"
```

- How to print line of the file /etc/passwd, that represents info about an user account "root" (see man 5 passwd)?

```
sed -n '/^root:/p' /etc/passwd
```

- How to print lines of the file /etc/passwd, that don't represent info about an user account "root"?

```
sed '/^root:/d' /etc/passwd
```

sed – text replacement

- Use the output of command `ls -l /etc`. How to replace the first character at line, which define the file type (d, -, l), by string (dir:, file:, link:).

- ```
ls -l /etc | \
sed 's/^d/dir:/' | \
sed 's/^-/file:/' | \
sed 's/^l/link:/'
```

- ```
ls -l /etc | \
sed 's/^d/dir:/;s/^-/file:/;s/^l/link:/'
```

- ```
cat types.sed
s/^d/dir:/
s/^-/file:/
s/^l/link:/
```

```
ls -l /etc | sed -f types.sed
```

# sed – text replacement

- Use the output of command `ls -l /`. How to replace all strings "root" by strings "ADMIN"?

```
ls -l / | sed 's/root/ADMIN/g'
```

- Use the output of command `ls -l /`. How to print only the second column by the command sed?

```
ls -l / | \
sed 's/^\([^]*\)[^]*$\) []*\([^]*\)[^]*$\)/\2/'
```

- Explain the behaviour of command `awk`.
- Example

```
ls -l /etc | \
awk '
 BEGIN { print "Regular files and directories" }
 /^-/ { f++; print f, $0; }
 /^d/ { d++; print d, $0 ;}
 END { print "Number of files:", f ;
 print "Number of directories:", d; }'
```



- How to print the first 10 lines of the file `/etc/passwd` by command `awk`?

```
cat -n /etc/passwd | awk 'NR<=10'
```

```
cat -n /etc/passwd | awk 'NR==1,NR==10'
```

```
cat -n /etc/passwd | \
awk '{ if (NR <= 10) { print $0;} }'
```

- How to print the last 10 lines of the file `/etc/passwd` by command `awk`?

```
cat -n /etc/passwd | \
awk "NR > $(($(wc -l </etc/passwd) - 10))"
```

- How to print lines containing string "root" from the file /etc/group by command awk?

```
awk '/root/' /etc/group
```

- How to print lines not containing string "root" from the file /etc/group by command awk?

```
awk '!/root/' /etc/group
```

- Use output of command `getent passwd`.
- How to print lines that represent account info about users whose names are Jan or Peter or Eliska?
  - Trivial solution

```
getent passwd | \
awk '/Jan/ ; /Peter/ ; /Eliska/'
```

- How to print lines that represent account info about users whose names are Jan or Peter or Eliska?
  - Correct solution

```
getent passwd | \
awk -F: \
'$5 ~ /^Jan / ; $5 ~ /^Peter ; / $5 ~ /^Eliska /'
```

```
cat names.awk
#!/usr/bin/awk
BEGIN { FS=":" }
$5 ~ /^Jan /
$5 ~ /^Peter /
$5 ~ /^Eliska /
```

```
getent passwd | awk -f names.awk
```

# awk – actions

- How to print the name and the size of every regular file in /etc/?

```
ls -l /etc | \
awk '/^-/ { printf("%s: %d B\n", $9,$5); }'
```

- How to print the user name and the path to home directory for every user that has account on this system?

Hint: Use command `getent passwd`.

```
getent passwd | \
awk 'BEGIN {FS=":"; OFS="..."}{print $1,$6}'
```

- How to print the user name of user that has an account on this system and has the longest name (the 5th column)?

```
getent passwd | \
awk -F: '{ l=length($5);
 if (l>m) { m=l; u=$5; } }
END{ print u; }'
```

- What is the meaning of the following awk functions?
  - `length(str)`  
... length of the string `str`,
  - `substr(str, from, len)`  
... the `len`-character substring of `str` starting at `from`,
  - `sub(ere, repl[, in])`  
... replace the first substring matching regular expression `ere` in the string `in`, by the string the `repl`,
  - `gsub(ere, repl[, in])`  
... replace all substring matching regular expression `ere`,
  - `split(str, arr[, fs])`  
... split the string `str` into the array `arr` according to the separators array `fs`.

- Assume that the file `m.txt` has the following contents:

|   |   |   |
|---|---|---|
| 1 | 2 | 3 |
| 8 | 7 | 9 |
| 3 | 7 | 2 |

- How to change the order of columns, such that the column  $i$  ( $0 \leq i < n$ ) will be swap with column  $(n - i)$  (reverse column order).

|   |   |   |
|---|---|---|
| 3 | 2 | 1 |
| 9 | 7 | 8 |
| 2 | 7 | 3 |

```
awk '{ for (i=NF; i>0; i--) printf("%d\t", $i);
 printf("\n"); }' m.txt
```

- Use output of command `ps -ef`. How to print info about processes which are running under the user `root`.

```
ps -ef | awk '/^ *root /'
```

```
ps -ef | nawk '{ if ($1 == "root") {print $0}; }'
```

- Use output of command `ps -eo user,rss,comm`. How to print the number of processes running under the effective identity "root" and the number of memory (RSS) allocated by these processes.
- Example of output

```
$> ps -eo user,rss,comm | awk -f ps1a.awk
root: nproc=85 rss=602264 KB
```

- Solution A

```
$> cat ps1a.awk
/^ *root /{
 nproc=nproc+1
 rss=rss+$2;
}
END {
 printf("root: nproc=%d rss=%d KB\n", nproc, rss)
}
```



- Solution B

```
$> cat ps1b.awk
BEGIN {
 user="root";
 nproc=0;
 rss=0;
}
{
 if ($1 == user)
 {
 nproc=nproc+1;
 rss=rss+$2;
 }
}
END {
 printf("%s: nproc=%d rss=%d KB\n", user, nproc, rss)
 printf("\n");
}
```

- Use the output from the command `ps -eo user,rss,comm`.  
How to print the following information about the processes of user whose name is given?

```
user_name: nproc=number_of_running_processes
rss=size_of_RSS_memory_alocatted_by_these_processes,
list=list_of_these_processes
```

- Example of output

```
$> ps -eo user,rss,comm | awk -f ps2.awk
root: nproc=83 rss=582952 KB
 proclist: sched, /usr/sbin/init, ...
```

## • Solution A

```
> cat ps2.awk
BEGIN {
 user="root";
 nproc=0;
 rss=0;
}
{
 if ($1 == user)
 {
 proclist[nproc]=$3;
 nproc=nproc+1;
 rss=rss+$2;
 }
}
END {
 printf("%s: nproc=%d rss=%d KB\n", user, nproc, rss)
 printf(" proclist: ");
 for (i=0; i<nproc; i++) {
 printf("%s, ", proclist[i]);
 };
 printf("\n");
}
```

## • Solution B

```
> cat ps3.awk
BEGIN {
 nproc=0; rss=0;
}
{
 if ($1 == user)
 {
 proclist[nproc]=$3;
 nproc=nproc+1;
 rss=rss+$2;
 }
}
END {
 printf("%s: nproc=%d rss=%d KB\n", user, nproc, rss)
 printf(" proclist: ");
 for (i=0; i<nproc; i++) {
 printf("%s, ", proclist[i]);
 };
 printf("\n");
}
```

```
$> ps -eo user,rss,comm | nawk -v user=root -f ps3.awk
```

# Homework

- Use the output from the command `ps -eo user,rss,comm`.  
How to print the following information for every user that is working on the server?

`user_name`, `nproc=number_of_running_processes`,  
`rss=size_of_RSS_memory_alocatted_by_these_processes`

- Example of output

```
$> ps -eo user,rss,comm | nawk -f ps4.awk
```

| User     | NProc | RSS (KB) |
|----------|-------|----------|
| -----    |       |          |
| kolensta | 2     | 28888    |
| root     | 94    | 689168   |
| trdlicka | 4     | 24448    |
| klamovik | 4     | 38112    |
| ...      |       |          |